

PulsePropagation

MMGNLSE Usage Manual

CPU and CUDA dense and CP-compressed propagation

Contents

1	Conventions	2
1.1	Power and photon field coordinates	2
2	A scalar single-mode propagation	3
3	Degrees of freedom and domains	5
4	Dispersion	5
4.1	Taylor coefficients	5
4.2	Sampled propagation constants	5
5	Lazy loss and gain arrays	6
6	Overlap tensors and polarization	6
6.1	Spatial overlap	6
6.2	Arbitrary component overlap	6
7	Materials and Raman responses	7
8	Building fiber properties	7
8.1	Analytic fibers	7
8.2	Custom index profiles	8
8.3	Reconstructing the spatial field	9
9	CP compression	9
10	CUDA execution	10
10.1	Selecting the backend	10
10.2	Capability matrix	10
10.3	Data movement and returned arrays	11
10.4	Reproducibility and CPU validation	12
10.5	Reusing a CP cache on a GPU worker	12
10.6	RunPod from SSH	13
11	Forward solves and saved planes	14
11.1	Accessing temporal and spectral fields	15
12	Observables	16
13	Adjoint solves	17
14	Reverse-mode autodiff adjoints	18
15	Stochastic ensembles	19
16	Explicit cache files and checksums	20
16.1	Fiber properties	20
16.2	CP decompositions	20
17	Worked examples	21
18	Troubleshooting	21

This manual covers the CPU and optional CUDA dense and CP-compressed MMGNLSE interface exposed by PulsePropagation.

1 Conventions

The package uses one set of units throughout the public interface:

Quantity	Unit
time	ps
angular-frequency detuning and ω_0	rad/ps
propagation distance and wavelength	m
optical power, $\text{abs2}(A)$	W
Taylor coefficient β_{k+1}	ps^k/m
overlap S	m^{-2}
nonlinear index n_2	m^2/W
loss α and gain	m^{-1}

A is a power-normalized complex field amplitude, so $\text{abs2}(A)$ is power in watts. The overlap remains an inverse area: the user supplies S , n_2 , and ω_0 separately.

1.1 Power and photon field coordinates

Power and photon units are two coordinate systems for the same field. Let $A_{\omega}[k]$ be the centered power-unit spectrum and define

$$w[k] = (Nt * dt) * 1e-12 / (\hbar * \omega_{\text{abs}}[k])$$

for positive absolute angular frequency $\omega_{\text{abs}}[k]$ in rad/s, with dt in ps and $\hbar = 1.054571817e-34$ J s. The weight is zero for a nonpositive absolute frequency. `photon_weights(grid,  $\omega_0$ )` returns these weights.

The forward photon-unit spectrum is

$$B_{\omega}[k] = \sqrt{w[k]} * A_{\omega}[k].$$

Consequently, $\text{abs2}(B_{\omega}[k])$ is the photon number in bin k , and summing it over bins, modes, and polarizations gives the total photon number. The inverse conversion divides by $\sqrt{w[k]}$; bins with zero weight remain zero.

Adjoints are cotangents, so their conversion is dual to the forward-field conversion. If λ_A is in power-field coordinates, then

$$\lambda_B[k] = \lambda_A[k] / \sqrt{w[k]}$$

for a positive-weight bin, and the inverse conversion multiplies by $\sqrt{w[k]}$. This preserves the differential pairing between a field perturbation and its adjoint. It is therefore important not to apply the forward-field scaling to an adjoint.

Both forward and adjoint field accessors default to `units=:power`. `units=:photon` first converts the centered spectral field. A temporal photon-unit field is then defined as

```
Bt = fft(ifftshift(Bomega, 1), 1)
```

It is the Fourier transform of the photon-unit spectrum, not a pointwise rescaling whose squared magnitude is a temporal photon flux.

Public modal fields always retain three axes:

```
time × mode × polarization
```

A saved forward or adjoint solution adds `z` last. A stochastic solution adds trajectory after `z`:

```
time × mode × polarization × z
time × mode × polarization × z × trajectory
```

Suppressed dimensions remain singleton dimensions. Linear polarization means two components ordered `x`, then `y`.

The public spectral convention is always centered:

```
Aomega = fftshift(ifft(At, 1), 1)
```

Thus frequency samples run from negative to positive detuning, including for odd grid sizes. `frequency_axis(grid)` returns angular-frequency detuning in rad/ps. `wavelength_axis(grid, omega0)` returns the corresponding absolute wavelengths in metres; bins with nonpositive absolute frequency are reported as `Inf`.

2 A scalar single-mode propagation

```
using PulsePropagation

grid = TimeGrid(1024, 10.0)           # 1024 samples over 10 ps
domain = MMGNLSEDomain(
    degrees_of_freedom(:time);
    time_grid=grid,
)

pulse_fwhm = 0.2                       # ps
peak_power = 2.0e3                     # W
initial = gaussian_pulse(
    domain, 1;
    peak_power=peak_power,
    fwhm=pulse_fwhm,
)

lambda0 = 1550e-9                     # m
c = 2.99792458e-4                     # m/ps
omega0 = 2pi * c / lambda0            # rad/ps
A_eff = 50e-12                        # m^2

parameters = MMGNLSEParameters(domain;
    length=0.1,                       # m
```

```

    alpha=0.0,
    gain=0.0,
    beta=TaylorBeta([0.0, 0.0, -0.0204]),
    S=1 / A_eff,
    n2=2.3e-20,
    omega0=omega0,
    raman=NoRaman(),
)

solution = solve_mmgnlse(initial, parameters, 1e-3; saveat=0.01)

At_final = get_temporal_field(solution)
Aomega_final = get_spectral_field(solution)
z_saved = solution.z

```

The initial field must have exactly the public shape implied by the domain and overlap. Here it is (1024, 1, 1). `gaussian_pulse` interprets `fwfm` as the intensity full width at half maximum and preserves `peak_power` as the total peak power across all launched components.

For a scalar multimode domain created with `degrees_of_freedom(:time, :space)`, provide a mode vector:

```

initial = gaussian_pulse(
    domain, nmodes;
    peak_power=2.0e3,
    fwhm=0.2,
    coefficients=[1.0, 0.5im, zeros(nmodes - 2)...],
)

```

For two linear polarizations, use a `mode × polarization` coefficient matrix. Rows correspond to modes and columns correspond to x and y:

```

coefficients = ComplexF64[
    1.0      0.25im
    0.5im    0.0
    0.0      exp(0.3im)
]

initial = gaussian_pulse(
    polarized_domain, 3;
    peak_power=2.0e3,
    fwhm=0.2,
    coefficients=coefficients,
)

```

The accepted coefficient formats are:

Domain	Required coefficients shape
scalar polarization	vector of length <code>nmodes</code>
linear x/y polarization	<code>nmodes × 2</code> matrix

A vector is rejected for a polarized domain, even when only x or y is to be launched; use a matrix with zeros in the other column. Coefficients may be real or complex, must all be finite, and must not all be zero.

The helper divides the complete vector or matrix by its Euclidean or Frobenius norm before applying the requested `peak_power`. Thus `peak_power` is the total peak power over all modes and polarizations, while relative complex amplitudes and phases are preserved.

With `coefficients=nothing`, the helper launches one component. `mode=1` and `polarization=1` are the defaults, and polarization indices one and two mean x and y. The `mode` and `polarization` keywords have no effect when explicit coefficients are supplied. Optional `time_offset` and `omega_offset` values are expressed in ps and rad/ps, respectively.

3 Degrees of freedom and domains

Create a nonempty selection from `:time`, `:space`, and `:polarization`:

```
dofs = degrees_of_freedom(:time, :space, :polarization)
domain = MMGNLSEDomain(dofs; time_grid=TimeGrid(2048, 20.0))
```

Names are canonicalized to time, space, polarization order. Calling `degrees_of_freedom()` with no arguments is an error.

Use only the dimensions the problem needs:

```
spatial_domain = MMGNLSEDomain(degrees_of_freedom(:space))
polarized_domain = MMGNLSEDomain(degrees_of_freedom(:polarization))
coupled_domain = MMGNLSEDomain(degrees_of_freedom(:space, :polarization))
```

When `:time` is active, a `TimeGrid` is required. When it is inactive, a grid must not be supplied, the retained time axis has length one, dispersion must be zero, and Raman must be disabled with `NoRaman()`.

4 Dispersion

4.1 Taylor coefficients

`TaylorBeta` stores coefficients about the carrier:

```
beta = TaylorBeta([beta0, beta1, beta2, beta3])
```

Its axes are coefficient order, mode, polarization. Omitted trailing axes are broadcast. For example:

```
coefficients = zeros(6, 3, 2) # beta0 through beta5, 3 modes, x/y
beta = TaylorBeta(coefficients)
@assert beta_order(beta) == 5
```

Propagation uses a moving frame: the `beta0` and `beta1` values of mode 1, polarization 1 define the common carrier phase and group-delay reference. Relative phase and walk-off of every other component are preserved.

4.2 Sampled propagation constants

Use `SampledBeta` for values sampled directly on the centered detuning grid:

```
omega = frequency_axis(grid) # rad/ps, negative to positive
values = beta_of_omega.omega # or frequency × mode × polarization
beta = SampledBeta(omega, values)
```

`omega` must be strictly increasing and must cover every solver detuning. `values` uses frequency, mode, polarization order and must contain exactly the grid's number of frequency samples.

5 Lazy loss and gain arrays

Loss and gain are power coefficients. Their field-amplitude contribution is

```
(gain - alpha) / 2.
```

Both accept a scalar or a positional array with axes

```
frequency × mode × polarization × z.
```

Rules for omitted dimensions are:

- A scalar broadcasts everywhere.
- A vector always means frequency dependence.
- An array occupies the axes shown above in order.
- Singleton dimensions explicitly skip an axis.
- Multiple `z` samples are uniformly placed over `[0, length]` and linearly interpolated.

Examples:

```
alpha_frequency = collect(range(0.01, 0.03; length=grid.nt))
alpha_mode = reshape([0.01, 0.02, 0.04], 1, 3)
gain_polarization = reshape([0.1, 0.08], 1, 1, 2)
gain_z = reshape([0.0, 0.2, 0.0], 1, 1, 1, 3)
```

Every supplied axis must either be singleton or match its target size, except the sampled `z` axis, whose `length` selects the number of uniform samples.

6 Overlap tensors and polarization

For a scalar single-mode problem, $S=1/A_{\text{eff}}$ is sufficient. A raw scalar or four-dimensional array passed to `MMGNLSEParameters` is converted to a `SpatialOverlap`.

6.1 Spatial overlap

```
overlap = SpatialOverlap(Smmmm) # Nm × Nm × Nm × Nm
```

`SpatialOverlap` stores only the spatial tensor. With two polarizations, the silica `x/y` tensor structure is applied analytically for both instantaneous and delayed nonlinearities. A one-polarization problem is its linear-`x` reduction.

6.2 Arbitrary component overlap

Use `ComponentOverlap` for a complete space-polarization tensor:

```
overlap = ComponentOverlap(Sfull)
```

`Sfull` has shape

```
Nm × 2 × Nm × 2 × Nm × 2 × Nm × 2.
```

No extra instantaneous polarization coefficients are applied. For delayed Raman, the supplied full component tensor defines the component couplings and the combined material response $h_a + h_b$ filters each $A3 \cdot \text{conj}(A4)$ beating term. Dense and component-CP forward and adjoint paths use the same rule.

7 Materials and Raman responses

Silica supplies Sellmeier dispersion, n_2 , and an anisotropic Raman response by default:

```
material = Silica()
material = Silica(n2=2.3e-20, raman=NoRaman())
material = Silica(raman=AgarwalRaman())
```

Available responses are:

```
NoRaman()

AgarwalRaman(
    fraction=0.18,
    tau1=0.0122,
    tau2=0.0320,
)

AnisotropicRaman(
    fraction=0.245,
    fa=0.75,
    fb=0.21,
    fc=0.04,
    tau1=0.0122,
    tau2=0.0320,
    taub=0.0960,
)
```

All time constants are in ps. Temperature is not part of these material responses. To inspect constructed causal kernels without supplying a sampled kernel yourself:

```
response = raman_response(AnisotropicRaman(), collect(0.0:0.001:0.5))
ha, hb = response.ha, response.hb
```

The Raman fraction is stored separately and is not folded into h_a or h_b .

8 Building fiber properties

8.1 Analytic fibers

```
grin = GRIN(
    lambda0=1550e-9,
    core_radius=31.25e-6,
    NA=0.275,
)
```

```
step = StepIndex(
    lambda0=1550e-9,
    core_radius=12.5e-6,
    NA=0.12,
)
```

The material index at every stencil wavelength comes from material dispersion. NA is held constant across the stencil. GRIN uses the standard parabolic n^2 profile; StepIndex uses weak-guidance scalar LP modes.

Compute modes, dispersion, and overlaps with:

```
properties = compute_fiber_properties(
    grin,
    Silica(),
    15;
    polarization=:linear,
    mode_basis=:HG,
    beta_order=5,
)
```

GRIN supports `mode_basis=:HG` and `:LP`. Step-index fibers require `:LP`. `polarization=:scalar` retains one polarization; `:linear` retains x/y while sharing the scalar spatial modes and overlap. `beta_order=5` is the default.

Useful accessors are:

```
xy = mode_grid(properties)           # (x=..., y=...)
modes = mode_fields(properties)      # y × x × mode
beta0 = propagation_constants(properties)
S = spatial_overlap(properties)
```

Mode fields satisfy discrete unit-power normalization. Construct propagation parameters directly from the result:

```
parameters = MMGNLSEParameters(domain;
    length=1.0,
    alpha=0.0,
    gain=0.0,
    beta=properties.beta,
    S=properties.S,
    n2=properties.n2,
    omega0=properties.omega0,
    raman=properties.raman,
)
```

8.2 Custom index profiles

A sampled custom profile uses a uniform Cartesian grid and stores its absolute index at `lambda0`:

```
profile = IndexProfile(x, y, nxy; lambda0=1030e-9)
```

`nxy` must have shape $(\text{length}(y), \text{length}(x))$. Away from `lambda0`, this form preserves its permittivity offset from material dispersion.

For explicit wavelength dependence, supply `n(x, y, lambda)` with all three arguments in metres:

```
index_model = (x, y, lambda) -> begin
  ncore = silica_refractive_index(lambda)
  x^2 + y^2 < (10e-6)^2 ? ncore : sqrt(ncore^2 - 0.12^2)
end

profile = IndexProfile(index_model; lambda0=1030e-9, x=x, y=y)
properties = compute_fiber_properties(
  profile, Silica(), 4;
  mode_basis=:FD,
  beta_order=5,
)
```

Custom modes are computed by the two-dimensional finite-difference Arpack solver. `evaluate_index_profile(profile, material, lambda)` evaluates any profile on its grid.

8.3 Reconstructing the spatial field

```
Axy = get_spatiotemporal_field(solution, properties)
Axy_omega = get_spatiotemporal_field(
  solution, properties; domain=:frequency)
```

The output axes are time/frequency, `x`, `y`, polarization. Stochastic output adds trajectory last and accepts `trajectory=:all` or a trajectory index. The reconstruction uses the carrier mode profiles stored in `properties`.

9 CP compression

Compress a spatial overlap with a deterministic complex CP decomposition:

```
S_cp = cp_compress(
  properties.S;
  error=1e-4,
  max_rank=1024,
  seed=20260722,
  checksum_fn=canonical_array_checksum,
  fftw_threads=1,
  blas_threads=1,
  npolarizations=2,
)
```

Passing `S_cp` as `S` in `MMGNLSEParameters` selects the compressed contraction automatically. The requested FFTW and BLAS thread counts are applied during compression and the previous process settings are restored afterward.

The fit uses complex factors and a nonconjugated reconstruction. With a fixed seed and settings it is deterministic. If the target error is not reached by `max_rank`, compression throws an error. A zero source tensor produces a rank-zero decomposition and propagation routes through the linear path.

Inspect or validate a result with:

```
rank = cp_rank(S_cp)
iszero_tensor = cp_iszero(S_cp)
S_approx = cp_reconstruct_mmgnlse(S_cp)
validate_cp_source(S_cp, properties.S.values;
    checksum_fn=canonical_array_checksum)
```

Compressing ComponentOverlap flattens the full space-polarization tensor and records a component layout automatically.

10 CUDA execution

CUDA is an optional dependency. Install it in the active project, then load CUDA once in the Julia process so the package extension is activated:

```
julia --project=. -e 'using Pkg; Pkg.add("CUDA"); Pkg.precompile()'
julia --project=. -e 'using CUDA; CUDA.versioninfo(); @assert CUDA.functional()'
```

The project accepts CUDA.jl major versions 5 and 6. A functional NVIDIA driver and a GPU assigned to the process are also required. Loading PulsePropagation without CUDA installed remains a CPU-only configuration.

10.1 Selecting the backend

The public solvers use symbols rather than exposing backend implementation types:

```
forward = solve_mmgnlse(
    initial,
    parameters,
    dz;
    method=RK4IP(),
    saveat=:steps,
    backend=:cuda,
    device=nothing,
    synchronize=true,
)
```

The same backend, device, and synchronize keywords are accepted by solve_adjoint and solve_mmgnlse_stochastic. The default is backend=:cpu. Use backend=:cuda for CUDA execution.

device=nothing uses CUDA.jl's current device. An integer selects a device; for example, device=0 selects the first visible GPU. Device indices refer to the process-visible devices after settings such as CUDA_VISIBLE_DEVICES have been applied.

Keep synchronize=true while validating or timing an end-to-end solve. Setting it to false skips the facade's explicit barrier. Required host copies still wait for their device work, so a returned solution is never a lazy or incomplete GPU result.

10.2 Capability matrix

The CUDA facade is designed to preserve the same physical parameters and array conventions as the CPU facade. "Regression-covered" below means the optional CUDA suite exercises the route when a functional device is available; on a machine without one, those tests are skipped.

Capability	backend=:cpu	backend=:cuda
Fixed-step RK4IP() forward solve and exact saves	yes	regression-covered
Fixed-step Tsit5() forward solve	yes	regression-covered
Dense SpatialOverlap, forward and adjoint	yes	regression-covered
Spatial CP decomposition, forward and adjoint	yes	regression-covered
Dense ComponentOverlap and component CP, forward and adjoint	yes	regression-covered
Scalar and linear x/y polarization, forward and adjoint	yes	regression-covered
Scalar and mode/polarization alpha and gain broadcasting	yes	regression-covered
Fully frequency-, mode-, polarization-, and z-varying alpha and gain	yes	regression-covered
TaylorBeta	yes	regression-covered
SampledBeta	yes	regression-covered
NoRaman, AgarwalRaman, and AnisotropicRaman, forward and adjoint	yes	regression-covered
Adjoint loss/gain signs and units=:power/:photon	yes	regression-covered
Default interaction-picture adjoint invocation	yes	regression-covered
Seeded input-noise stochastic ensemble	yes	regression-covered
Reverse-mode Zygote autodiff adjoint	checkpointed	no
Fiber-mode construction and cp_compress	host	host

CUDA forward propagation supports fixed-step `method=RK4IP()` and fixed-step `method=Tsit5()`. For a CUDA adjoint, use the public default `Vern9()` method selection rather than passing an alternative integration method. The returned adjoint still follows the public centered-spectral terminal convention and has ascending saved `z` coordinates.

Fiber-mode construction, CP fitting, JLD2/TOML cache validation, observable evaluation, and result serialization run on the host. A verified CP decomposition can therefore be prepared once on a CPU machine and reused by CUDA solves.

10.3 Data movement and returned arrays

Pass ordinary Julia arrays and the same `MMGNLSEParameters` used for CPU execution. With `backend=:cuda`, active fields, operators, overlap factors, and solver workspaces are copied to and retained on the selected GPU during the numerical solve. Public solution data is copied back before return:

```
@assert forward.fields isa Array
@assert get_spectral_field(forward) isa Array
```

Forward, adjoint, and stochastic results therefore have exactly the host array axes documented above. Do not wrap the public initial field or overlap in a device-array type. Saved trajectories can be the dominant host transfer and memory cost; use `saveat=nothing` when only endpoints are needed and `saveat=:`

steps when the adjoint or a longitudinal contraction requires every landed plane.

10.4 Reproducibility and CPU validation

Use explicit seeds for both CP fitting and stochastic input noise:

```
S_cp = cp_compress(
    properties.S;
    error=1e-3,
    seed=20260724,
    checksum_fn=canonical_array_checksum,
)

ensemble = solve_mmgnlse_stochastic(
    initial, parameters, dz, 100;
    noise=1.0,
    seed=13792,
    backend=:cuda,
)
```

The stochastic seed fixes the host-generated input perturbations, including independent polarization samples. It does not make floating-point reductions on different GPU models bitwise identical. Compare CPU and GPU calculations with tolerances and physics-level observables, not byte equality.

For a new parameter regime, first run a reduced problem on both backends:

```
cpu = solve_mmgnlse(
    initial, parameters, dz;
    method=RK4IP(), saveat=nothing, backend=:cpu)
gpu = solve_mmgnlse(
    initial, parameters, dz;
    method=RK4IP(), saveat=nothing, backend=:cuda,
    synchronize=true)

relative_l2 = norm(
    get_spectral_field(gpu) .- get_spectral_field(cpu)) /
    max(norm(get_spectral_field(cpu)), eps(Float64))
@assert relative_l2 <= 2e-7
```

The $2e-7$ forward threshold is a practical smoke-test value for the reduced worked example, not a universal error guarantee. Its example adjoint uses $2e-5$. Tighten or relax tolerances only after recording the grid, step size, precision, GPU model, CUDA versions, CP checksum, and the corresponding CPU convergence study.

For adjoint validation, form `lambda_L` from the same observable on each backend's forward result, compare the input-plane spectral adjoints, and also check a centered finite-difference directional derivative. Save the forward with `saveat=:steps` during this check. `solve_adjoint` returns its trajectory in ascending `z` order even though it integrates from the output plane toward the input.

10.5 Reusing a CP cache on a GPU worker

CP caches are portable host files and retain the source-tensor checksum:

```
S_dense = properties.S.values
S_cp = load_cp_decomposition(
```

```

    "/workspace/cache/step_index_cp.jld2",
    "/workspace/cache/step_index_cp.toml";
    source_S=S_dense,
)

parameters = MMGNLSEParameters(domain;
    length=length_m,
    beta=properties.beta,
    S=S_cp,
    n2=properties.n2,
    omega0=properties.omega0,
    raman=properties.raman,
)

```

Keep both cache files together. A partial pair, file corruption, or a different dense overlap is rejected before GPU work begins.

10.6 RunPod from SSH

Choose a CUDA-capable pod, expose SSH, and attach a persistent volume at `/workspace` if caches and results must survive pod termination. After connecting:

```

nvidia-smi

# Install Julia when the pod image does not already provide it.
curl -fsSL https://install.julialang.org | sh -s -- -y
export PATH="$HOME/.juliaup/bin:$PATH"
juliaup add 1.12
juliaup default 1.12

cd /workspace
git clone https://github.com/nrivera494/FastDifferentiableMMNLO.git
cd FastDifferentiableMMNLO

julia --project=. -e 'using Pkg; Pkg.instantiate(); Pkg.add("CUDA"); Pkg.precompile()'
julia --project=. -e 'using CUDA; CUDA.versioninfo(); @assert CUDA.functional()'
julia --project=. test/test_mmgnlse_cuda_optional.jl

```

Run the reduced workflow first. Its default performs CPU/GPU forward and adjoint checks and writes both its CP cache and result file under `/workspace`:

```

mkdir -p /workspace/pulseprop-cache /workspace/pulseprop-results

PULSEPROP_BACKEND=cuda \
PULSEPROP_VALIDATE_CPU=true \
PULSEPROP_CACHE_DIR=/workspace/pulseprop-cache \
PULSEPROP_OUTPUT_DIR=/workspace/pulseprop-results \
julia --project=. examples/new_api/step_index_squeezing_gpu.jl

```

Then run the notebook-scale grid and propagation length:

```

PULSEPROP_BACKEND=cuda \
PULSEPROP_FULL_EXAMPLES=true \
PULSEPROP_VALIDATE_CPU=false \
PULSEPROP_NTRAJ=100 \

```

```
PULSEPROP_CACHE_DIR=/workspace/pulseprop-cache \
PULSEPROP_OUTPUT_DIR=/workspace/pulseprop-results \
julia --project=. examples/new_api/step_index_squeezing_gpu.jl
```

The full branch keeps the notebook’s numerical settings: $N_t=2048$, a 10 ps window, three scalar LP modes, 2.86 m length, dense-forward $dz=2.86/4000$ m, CP-forward $dz=1e-3$ m, $saveat=0.1$ m for both forwards, adjoint $dz=2.86/4000$ m, CP target error $1e-3$, and 100 stochastic trajectories. CPU/GPU parity, when enabled, repeats those settings rather than substituting a cheaper reference.

Useful controls are:

Environment variable	Meaning	Default
PULSEPROP_BACKEND	cpu or cuda	cuda
PULSEPROP_FULL_EXAMPLES	notebook-scale settings	false
PULSEPROP_CUDA_DEVICE	visible integer device index	current device
PULSEPROP_CUDA_SYNCHRONIZE	explicit CUDA synchronization	true
PULSEPROP_VALIDATE_CPU	CPU reference checks at the selected resolution	true for reduced CUDA
PULSEPROP_FORWARD_RTOL	reduced forward relative-L2 limit	$2e-7$
PULSEPROP_ADJOINT_RTOL	reduced adjoint relative-L2 limit	$2e-5$
PULSEPROP_CP_SEED	deterministic CP fit seed	20260724
PULSEPROP_NOISE_SEED	deterministic input-noise seed	13792
PULSEPROP_NTRAJ	stochastic trajectory count	8 reduced, 100 full
PULSEPROP_CACHE_DIR	paired CP cache location	repository output cache
PULSEPROP_OUTPUT_DIR	JLD2 result location	repository output

The result JLD2 contains host arrays for the final dense and CP temporal and spectral fields, their relative error, the input adjoint, stochastic filtered-photon samples, saved z coordinates, checksums, seeds, settings, and any CPU/GPU validation errors. Copy that file and both CP cache files off ephemeral storage before stopping the pod.

11 Forward solves and saved planes

```
solution = solve_mmgnlse(
    initial,
    parameters,
    dz;
```

```

    method=RK4IP(),
    saveat=nothing,
)

```

RK4IP() is the default fixed-step method. Tsit5() is also available and is run with adaptivity disabled at the requested dz.

Saving rules are:

- saveat=nothing: save only the beginning and end.
- saveat=:steps: save every landed integration step.
- saveat=spacing: save at that distance interval.
- saveat=[z1, z2, ...]: save at the requested positions.

The exact endpoint is always appended, and steps are shortened as necessary to land on every requested plane. solution.z is the authoritative list of saved coordinates.

11.1 Accessing temporal and spectral fields

Use the accessors rather than relying on the stored representation:

```

At = get_temporal_field(solution; z=:final, units=:power)
Aomega = get_spectral_field(solution; z=:all, units=:photon)
Bt = get_temporal_field(solution; z=:all, units=:photon)

lambda_t = get_temporal_field(adjoint; z=:all, units=:power)
lambda_omega_0 = get_spectral_field(
    adjoint; z=:initial, units=:photon)

At_ensemble = get_temporal_field(
    ensemble; z=:final, trajectory=:all, units=:power)

```

The public defaults are:

Object	Default z	Default units	Other selector
forward solution	:final	:power	none
adjoint solution	:initial	:power	none
stochastic ensemble	:final	:power	trajectory=:all

These defaults also apply to calls without keywords, so existing calls such as get_spectral_field(solution) continue to work and return the final power-unit field. For either accessor, z accepts :initial, :final, :all, a saved-plane integer index, or a real coordinate. A real coordinate selects the nearest saved plane and does not interpolate. z=:all retains the saved-z axis. A stochastic trajectory selector accepts :all or one integer trajectory index.

get_spectral_field returns the centered negative-to-positive spectrum

```

Aomega = fftshift(iffshift(At, 1), 1)

```

with no additional public prefactor. get_temporal_field applies the inverse operation after any requested unit conversion. Forward solutions are stored as power-unit temporal fields. An adjoint may

be stored in the units requested from `solve_adjoint`, but its accessors still default to `units=:power` and perform the needed conversion. The `units=:photon` definitions, including the dual conversion for adjoints, are given earlier under “Power and photon field coordinates.”

12 Observables

All spectral observables use centered ordering. The most common constructor defaults are:

Constructor	Required argument	Defaults	Meaning with defaults
SpectralPhotonNumber	filter	modes=:all	filtered photon number over all modes and polarizations
FilterEnergy	filter	domain=:frequency, modes=:all	filtered squared-field sum in frequency coordinates
BinEnergy	index	domain=:frequency	squared field at one frequency bin; omitted mode or polarization indices default to one
TemporalMoment	none	order=1, center=0.0, normalized=true, modes=:all	energy-normalized temporal centroid
SpectralMoment	none	order=1, center=0.0, normalized=true, modes=:all	energy-normalized mean angular-frequency detuning
TemporalCentroid	none	modes=:all	alias for the default temporal moment
SpectralCentroid	none	modes=:all	alias for the default spectral moment

For `SpectralPhotonNumber`, write `filter=nothing` explicitly to include every spectral bin, or use a scalar filter such as `filter=1.0`. Filters may be a real scalar or a lazy array with axes frequency or time, mode, polarization; singleton dimensions broadcast. A filter weights `abs2(field)` in the objective. For an amplitude transfer function `H`, therefore supply `abs2.(H)` as the filter. `modes` selects spatial modes; all polarizations are retained.

For the variance (squared RMS width), use `order=2`, `center=:centroid`, and `normalized=true`; take `sqrt(value(...))` for the RMS width. In contrast, `order=1`, `center=:centroid` is the first central moment and is zero, not the centroid. Setting `normalized=false` returns the unnormalized field-weighted moment.

Common constructors are:

```
photon_observable = SpectralPhotonNumber(filter=filter, modes=:all)
energy_observable = FilterEnergy(
    filter=filter, domain=:frequency, modes=:all)
bin_observable = BinEnergy(index=(frequency_index, mode, polarization))
temporal_centroid = TemporalCentroid(modes=:all)
spectral_centroid = SpectralCentroid(modes=:all)
temporal_width = TemporalMoment(
    order=2, center=:centroid, normalized=true, modes=:all)
```

```
spectral_width = SpectralMoment(
    order=2, center=:centroid, normalized=true, modes=:all)
ratio = RatioObservable(photon_observable, energy_observable)
```

Evaluate an observable and obtain its centered spectral terminal condition:

```
J = value(photon_observable, solution)
lambda_L = terminal_condition(photon_observable, solution)
```

`value(observable, solution)` evaluates the final saved plane by default. `value(observable, ensemble)` returns one value per stochastic trajectory; statistical summaries such as mean and var are ordinary postprocessing.

`lambda_L` is the Wirtinger derivative $dJ/d\text{conj}(A\omega(L))$ and can be passed directly to `solve_adjoint`.

Convenience diagnostics include:

```
photons = photon_spectrum(solution)
N = photon_number(solution, filter; modes=:all)
energy_nJ = modal_energy(solution; sum_polarizations=false)
weights = photon_weights(grid, parameters.omega0)
```

`photon_spectrum(...; normalization=:bin)` returns photons per bin; `normalization=:density` returns density per angular-frequency sampling interval. Nonpositive absolute-frequency bins have zero photon weight. Photon-number diagnostics require the `:time` degree of freedom because the bin normalization depends on the temporal window.

13 Adjoint solves

Given a saved forward solution:

```
lambda_L = terminal_condition(observable, forward)

adjoint = solve_adjoint(
    lambda_L,
    parameters,
    dz_adj;
    forward_solution=forward,
    method=Vern9(),
    units=:power,
    saveat=:steps,
)
```

Alternatively, ask the routine to generate its own forward trajectory:

```
adjoint = solve_adjoint(
    lambda_L,
    parameters,
    dz_adj;
    initial_field=initial,
    dz_forward=dz,
)
```

Supply exactly one of `forward_solution` and `initial_field`. When an initial field is supplied, `dz_forward` is required. A sparsely saved forward solution is replayed at its integration steps before the adjoint solve.

The adjoint is integrated in the interaction picture with fixed-step `Vern9()` applied only to the nonlinear VJP. Dispersion and longitudinally varying gain/loss use their exact diagonal propagators, matching the forward RK4IP linear convention. Returned `z` coordinates are ascending even though integration proceeds backward. The input `lambda_L` is always expressed in power-field coordinates. `units=:photon` converts the saved returned gradient to $dJ/d\text{conj}(B)$ and zeros nonphysical bins; `units=:power` retains power-field coordinates.

Dense and CP paths support loss, gain, spatial or complete component tensors, linear x/y polarization, `NoRaman()`, `AgarwalRaman()`, and `AnisotropicRaman()`.

14 Reverse-mode autodiff adjoints

`get_autodiff_adjoint` computes an independent reverse-mode derivative through the discrete forward solver:

```
lambda_omega_0 = get_autodiff_adjoint(
    observable,
    initial,
    parameters,
    dz;
    method=RK4IP(),
    units=:power,
    checkpointing=:auto,
)
```

The function performs its own forward solve, constructs the observable's terminal condition, and differentiates the fixed-step RK4IP propagation with `Zygote`. It returns the centered spectral cotangent at $z=0$ with shape `(frequency, mode, polarization)`. It does not return an adjoint solution or a saved trajectory.

This result is the exact reverse derivative of the *discretized* RK4IP propagation. It is intended to agree, as the step size converges, with

```
get_spectral_field(
    solve_adjoint(...);
    z=:initial,
    units=units,
)
```

where `solve_adjoint` integrates the continuous interaction-picture adjoint. The two routes are therefore a useful regression comparison but need not be bitwise equal at finite `dz`. Use `solve_adjoint` when a longitudinal adjoint trajectory is required.

The checkpointing choices are:

checkpointing	Behavior
:auto	blocks of 10 RK4IP steps

Continued on next page

checkpointing	Behavior
positive integer <i>n</i>	at most <i>n</i> RK4IP steps per block
:none	retain the full reverse-mode tape

Checkpointing stores only block boundaries and recomputes each block during the reverse pass. It therefore trades roughly one additional forward trajectory of computation for a much smaller active tape. A smaller block uses less memory and more boundaries; a larger block uses more memory.

The current autodiff route is CPU-only, uses Zygote, and accepts only `method=RK4IP()`. `units=:power` is the default; `units=:photon` applies the dual adjoint conversion described above and zeros nonpositive absolute-frequency bins. CUDA forward and continuous-adjoint solves remain available through `backend=:cuda`, but GPU reverse-mode autodiff is not currently exposed.

15 Stochastic ensembles

```
ensemble = solve_mmgnlse_stochastic(
    initial,
    parameters,
    dz,
    100;
    noise=1.0,
    seed=1234,
    saveat=:steps,
)
```

noise is an input photon-number Fano factor:

- A scalar is uniform across frequency, mode, and polarization.
- A vector is frequency-dependent.
- Positional arrays use frequency, mode, polarization order and may use singleton dimensions.

Stochastic photon-noise ensembles require the `:time` degree of freedom.

Each frequency-mode-polarization component is independent and sampled as

```
delta_B = sqrt(F)/2 * (xi1 + i*xi2),
```

where `xi1` and `xi2` are independent unit normal variables. The `x` and `y` noises are independent with the same Fano factor when that axis is broadcast. Supplying a fixed `seed` makes the ensemble deterministic. `seed` and `rng` cannot be supplied together.

Access one or all trajectories with:

```
At_one = get_temporal_field(ensemble; trajectory=1)
Aomega_all = get_spectral_field(ensemble; trajectory=:all)
```

This routine samples initial fluctuations only. It does not add distributed Raman or gain Langevin sources. Consequently, complete QSA squeezing and timing-jitter curves that require longitudinal source contractions need an additional calculation over the saved forward and adjoint `z` coordinates. There is no temperature keyword for this input-noise sampler; temperature belongs with the correlation model used for a distributed Langevin source.

16 Explicit cache files and checksums

Fiber properties and CP decompositions use paired JLD2 data and TOML metadata files. Paths are explicit, parent directories must already exist, and existing files are protected unless `overwrite=true` is supplied.

16.1 Fiber properties

```
save_fiber_properties(  
    "cache/fiber.jld2",  
    "cache/fiber.toml",  
    properties,  
)  
  
checksum = fiber_properties_checksum(properties)  
  
loaded_properties = load_fiber_properties(  
    "cache/fiber.jld2",  
    "cache/fiber.toml";  
    expected_checksum=checksum,  
    source_S=properties.S.values,  
)
```

16.2 CP decompositions

Attach source identity during compression or saving:

```
S_dense = properties.S.values  
S_cp = cp_compress(  
    S_dense;  
    error=1e-4,  
    max_rank=1024,  
    checksum_fn=canonical_array_checksum,  
)  
  
save_cp_decomposition(  
    "cache/overlap_cp.jld2",  
    "cache/overlap_cp.toml",  
    S_cp;  
    source_S=S_dense,  
)  
  
loaded_cp = load_cp_decomposition(  
    "cache/overlap_cp.jld2",  
    "cache/overlap_cp.toml";  
    source_S=S_dense,  
)
```

Loading validates the JLD2 SHA-256 checksum, individual arrays, logical payload metadata, and any supplied expected/source checksum. A changed source tensor is rejected.

17 Worked examples

Five executable workflows are provided:

- [birefringent_mm_soliton.jl](#)
- [gif625_analytic.jl](#)
- [squeezing_spectral_filter.jl](#)
- [soliton_jitter.jl](#)
- [step_index_squeezing_gpu.jl](#)

They use reduced settings by default. Select the notebook-scale grids, pulse parameters, fiber lengths, and step sizes with:

```
PULSEPROP_FULL_EXAMPLES=true julia --project=. \
  examples/new_api/birefringent_mm_soliton.jl
```

The squeezing and jitter workflows exercise the forward, adjoint, filter, and input-noise pieces. Their distributed Langevin-source contribution is not included, as described in the stochastic section.

The step-index squeezing script is the executable reduced/full translation of [step_index_squeezing.ipynb](#). It defaults to CUDA, caches its CP decomposition, uses deterministic input noise, optionally validates against CPU, and writes a portable JLD2 result. See the RunPod section for complete commands.

18 Troubleshooting

Symptom	Check
A field-shape error	Retain time, mode, polarization axes, including singleton dimensions.
A sampled-dispersion error	Use <code>frequency_axis(grid)</code> in rad/ps and supply exactly <code>Nt</code> increasing samples.
A no-time parameter error	Use zero dispersion, <code>NoRaman()</code> , no <code>TimeGrid</code> , and an initial time axis of length one.
A loss/gain axis error	Use frequency, mode, polarization, z order; add singleton dimensions to reach a later axis.
A CP compression failure	Relax error, increase <code>max_rank</code> , or increase <code>maxiter</code> ; missed targets never return silently.
A cache checksum error	Load with the same dense source tensor and an unmodified JLD2/TOML pair.
Unexpected spectral order	Public spectra and axes are already centered; do not apply an additional shift.
Zygote autodiff exhausts memory	Use <code>get_autodiff_adjoint(...; checkpointing=:auto)</code> or a smaller positive block size; use <code>solve_adjoint</code> when a full trajectory or lower memory is more important than differentiating the discrete solver.
A full example is slow	Run without <code>PULSEPROP_FULL_EXAMPLES</code> first; the full grids and trajectory counts are intentionally computationally intensive.

Continued on next page

Symptom	Check
CUDA.jl loads but no GPU is usable	Check <code>nvidia-smi</code> , pod GPU assignment, <code>CUDA_VISIBLE_DEVICES</code> , then run using <code>CUDA; CUDA.versioninfo(); CUDA.functional()</code> .
A CUDA method error	Use fixed-step <code>RK4IP()</code> or fixed-step <code>Tsit5()</code> for forward propagation and fixed-step <code>Vern9()</code> for the adjoint.
CPU/GPU validation misses tolerance	First reduce <code>dz</code> , confirm the same CP checksum and saved planes, keep <code>synchronize=true</code> , and compare observables as well as relative L2 error.
A RunPod result disappears	Store output and both cache files on a mounted persistent volume before terminating the pod.